

Legion: Expressing locality and independence with logical regions

Michael Bauer, Sean Treichler, Elliott Slaughter, Alex Aiken

Stanford University

April 24, 2025

Intuition

From (Restricted) Boltzmann machine to Helmholtz Machine

Boltzmann machine's energy function (a column vector $\{0, 1\}$)

$$p(\mathbf{x}) = \frac{\exp(-E(\mathbf{x}))}{Z}$$


$$E(\mathbf{x}) = -\mathbf{x}^T U \mathbf{x} - \mathbf{b}^T \mathbf{x} = -\sum_{i,j} w_{i,j} x_i x_j - \sum_i w_i x_i$$

Input can be divided into two parts:

$$\mathbf{x} = (\mathbf{v}, \mathbf{h})$$

Restricted Boltzmann Machine

visible variables and hidden variables are fully connected



Restricted Boltzmann Machine

$$p(\mathbf{v}, \mathbf{h}) = \frac{\exp(-E(\mathbf{v}, \mathbf{h}))}{Z}$$

$$E(\mathbf{v}, \mathbf{h}) = -\mathbf{v}^T W \mathbf{h} - b^T \mathbf{v} - c^T \mathbf{h}$$

Then we can get the conditional distribution for hidden variables given visible variables (encoder):

$$p(\mathbf{h} | \mathbf{v}) = \prod_{j=1}^{n_h} \sigma\left((2h - 1) * \left(c + W^T \mathbf{v}\right)\right)_j$$

Restricted Boltzmann Machine

Decoder:

$$p(\mathbf{v} \mid \mathbf{h}) = \prod_{i=1}^{n_v} \sigma\left((2v_i - 1) * (b_i + W_i^T \mathbf{h})\right)$$



Learning with Contrastive Divergence.

Helmholtz Machine

Helmholtz Machine is a hierarchical full-connected feedback NN.

RBM can be viewed as a Helmholtz machine with only **one layer of hidden units** and an **implicit posterior distribution**.

Helmholtz Machine



Divide generative model and recognition model

Why it is called Helmholtz Machine? Two models are trained iteratively.

MLE, theta is model parameters:

$$\theta = \arg \max(\theta) \sum_{i=1}^n \log p(\mathbf{v}_i | \theta)$$

$$\theta = \arg \max(\theta) \sum_{i=1}^n \log p(\mathbf{v}_i, \mathbf{h} | \theta)$$

Real posterior distribution is intractable

$$p(\mathbf{h} \mid \mathbf{v}, \theta) = \frac{p(\mathbf{h}, \mathbf{v} \mid \theta)}{p(\mathbf{v} \mid \theta)}$$

Approximation with Q:

$$Q(\mathbf{h} \mid \mathbf{v}) = \prod_{i=1}^Q (h_i \mid \mathbf{v})$$

ELBO (Evidence Lower Bound) from variational inference:

$$\log p(\mathbf{v} \mid \theta) = E_Q[\log p(\mathbf{v}, \mathbf{h} \mid \theta) - \log Q(\mathbf{h} \mid \mathbf{v})] + D_{KL}(Q(\mathbf{h} \mid \mathbf{v}) \parallel p(\mathbf{h} \mid \mathbf{v}, \theta))$$

Real posterior distribution is intractable

$$D_{KL}(Q(\mathbf{h} \mid \mathbf{v}) \parallel p(\mathbf{h} \mid \mathbf{v}, \theta)) > 0$$

Therefore, the ELBO is a lower bound of the log-likelihood.

$$L(\theta, Q) = E_Q[\log p(\mathbf{v}, \mathbf{h} \mid \theta) - \log Q(\mathbf{h} \mid \mathbf{v})]$$

That is, $\log p(\mathbf{v} \mid \theta) \geq L(\theta, Q)$.

Introduce free energy

$$F(\mathbf{v}, \theta) = -E_Q[\log p(\mathbf{h}, \mathbf{v} \mid \theta)] - H(Q(\mathbf{h} \mid \mathbf{v}))$$

The objective is to minimize the free energy.

Now we have two models: $p(\mathbf{v}, \mathbf{h} \mid \theta)$ and $Q(\mathbf{h} \mid \mathbf{v}; \varphi)$

We can use the free energy to optimize the parameters.

$$\theta = \arg \min(\theta) F(\mathbf{v}, \theta)$$

$$\varphi = \arg \min(\varphi) F(\mathbf{v}, \theta)$$

Bi-level optimization

Wake-sleep algorithm The wake phase is to minimize the free energy, and the sleep phase is to maximize the free energy.

Wake phase: target is minimize $F(\mathbf{v}, \theta)$, i.e., maximize $L(\theta, Q)$. use real data pass Q to get posterior distribution. then update θ .

Sleep phase: generate data with $p(\mathbf{v}, \mathbf{h} \mid \theta)$. pass generated data to Q to get posterior distribution. then update ϕ .

Modern development


Variational Autoencoders (VAEs)

Generative Adversarial Networks (GANs)

Flow-based models

Diffusion models

Variational Autoencoders (VAEs)



Works on LLM

KV-cache compression via importance evaluation

Block Attention